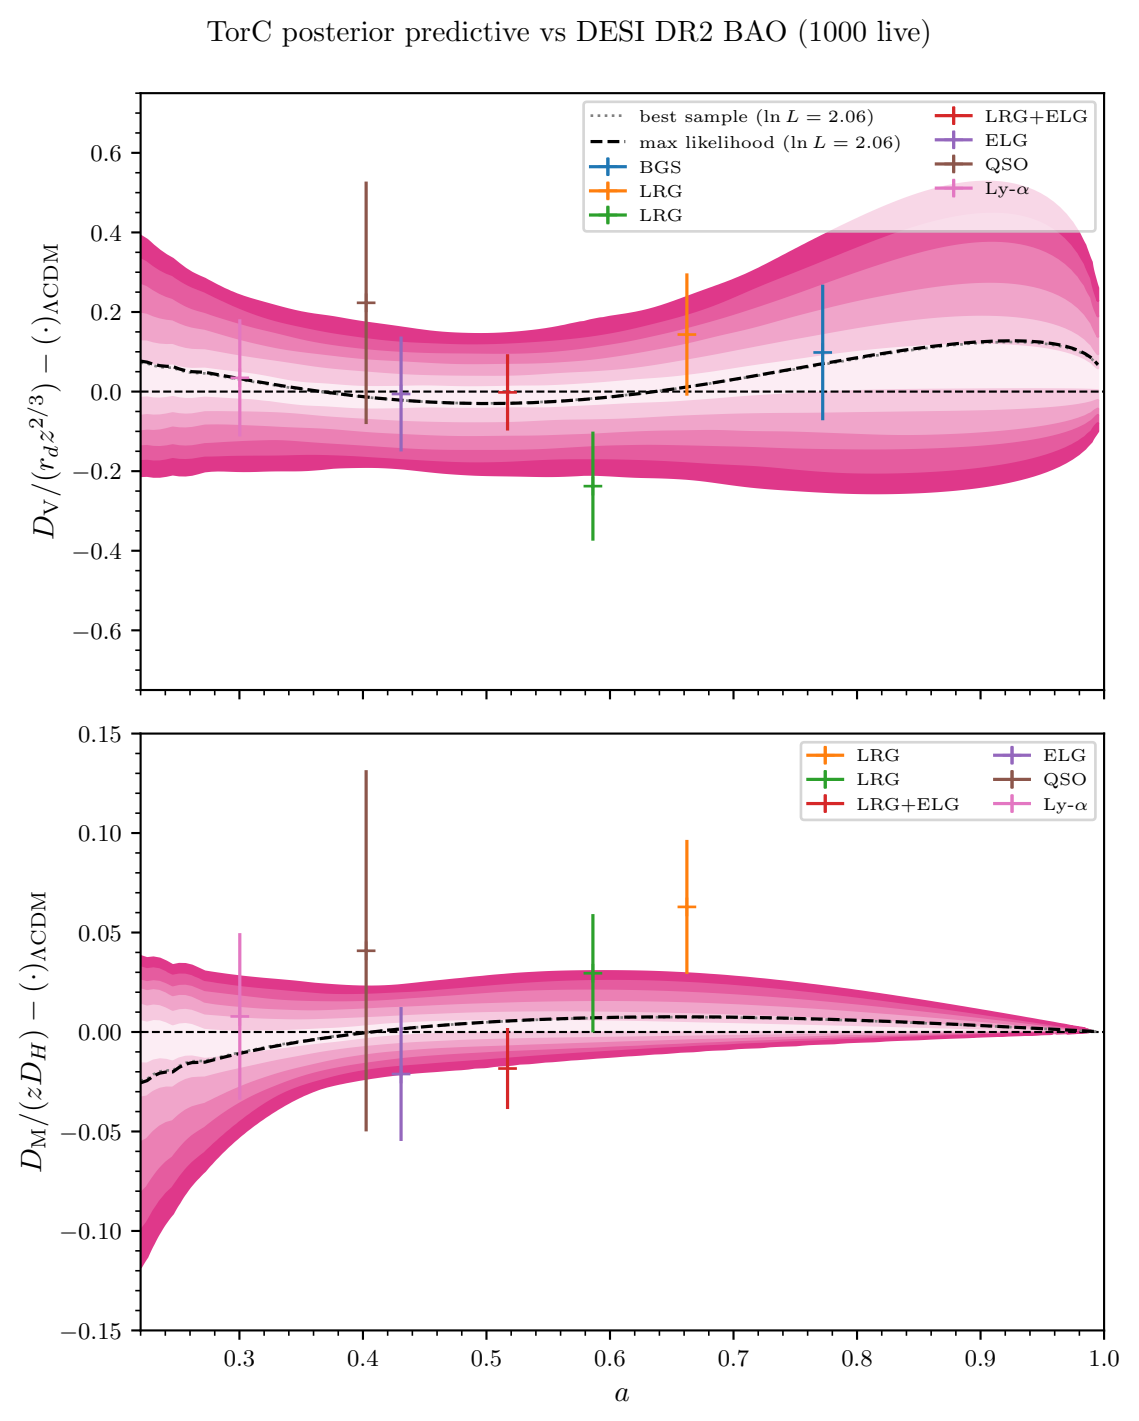


max-likelihood curve: Nelder-Mead on $\ln L$, started from the best of 13500 dead points, bounded to the prior box.
$\ln L$: 2.058 (best sample) → 2.062 (optimised), gain +0.004 in 662 evaluations.
ΛCDM reference curve is fitted the same way (χ^2, Nelder-Mead).